

This notebook is for Experiment 3: K-Means After PCA

The aims of this notebook is to achieve the following:

- Test with different numbers of principal components (2, 5, 10, 15, 20)
- Analyze trade-off between dimensionality and clustering quality
- Compare reconstruction error vs clustering performance

For ease of visualization and code readability, all codes previously made (K-Means and PCA) will be rewritten in the notebook.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_breast_cancer
from sklearn.metrics import confusion_matrix
import seaborn as sb

# Load the breast cancer diagnostic dataset
dataset = load_breast_cancer()
features = dataset.data
target = dataset.target
```

# Metrics:

```
def mse(y_true, y_pred):
    return np.mean((y_true - y_pred)**2)

def silhouette_score(X, labels):
    clusters = np.unique(labels)
    samples = X.shape[0]

    # Compute distances matrix
    distances = np.zeros((samples, samples))
    for i in range(samples):
        for j in range(samples):
            distances[i, j] = np.linalg.norm(X[i] - X[j])

    # Compute scores
    scores = np.zeros(samples)
    for i in range(samples):
```

```
cluster = labels[i]
# Cohesion
same_samples = np.where(labels == cluster)[0]
same_samples = same_samples[same_samples != i]

if len(same_samples) == 0:
    cohesion = 0
else:
    cohesion = np.mean(distances[i, same_samples])

# Separation
separation = np.inf
for j in clusters:
    if j != cluster:
        separation_j = np.mean(distances[i, labels == j])
        if separation_j < separation:
            separation = separation_j

scores[i] = (separation - cohesion) / max(separation, cohesion)

return np.mean(scores)

def davies_bouldin_index(X, labels):
    clusters = np.unique(labels)
    samples = X.shape[0]

    centroids = []
    dispersion = []

    # Intra-Cluster
    for cluster in clusters:
        cluster_points = X[labels == cluster]
        centroid = np.mean(cluster_points, axis=0)
        centroids.append(centroid)
        dispersion.append(np.mean(np.linalg.norm(cluster_points - c

dispersion = np.array(dispersion)
centroids = np.array(centroids)

# Inter-Cluster
distances = np.linalg.norm(centroids[:, np.newaxis] - centroids

max_ratios = []
for i in range(len(clusters)):
    current_index_ratio = []
    for j in range(len(clusters)):
        if i != j:
            ratio = (dispersion[i]+dispersion[j]) / distances[i, j]
```

```
        current_index_ratio.append(ratio)
    max_ratios.append(np.max(current_index_ratio))

    return np.mean(max_ratios)

def calinski_harabasz_index(X, labels):
    clusters = np.unique(labels)
    samples = X.shape[0]

    mean_global = np.mean(X, axis=0)

    # Between-Cluster Variance
    between_cluster_sum_of_squares = 0
    within_cluster_sum_of_squares = 0

    for cluster in clusters:
        cluster_points = X[labels == cluster]
        mean_cluster = np.mean(cluster_points, axis=0)

        nSamples = cluster_points.shape[0]
        within_cluster_sum_of_squares += np.sum((cluster_points - mean_global)**2)
        between_cluster_sum_of_squares += nSamples * np.sum((cluster_points - mean_global)**2)

    index = ((between_cluster_sum_of_squares / (len(clusters) - 1)) -
              within_cluster_sum_of_squares / samples) / within_cluster_sum_of_squares

    return within_cluster_sum_of_squares, index

def adjusted_rand_index(labels_true, labels_pred):
    TP = 0
    TN = 0
    FP = 0
    FN = 0

    for i in range(0, len(labels_pred)-1):
        for j in range(i+1, len(labels_pred)):
            if labels_pred[i] == labels_pred[j]:
                # Positive Pair
                if labels_true[i] == labels_true[j]:
                    # True
                    TP += 1
                else: #False
                    FP += 1
            else: # Negative Pair
                if labels_true[i] == labels_true[j]:
                    # False
                    FN += 1
                else: # True
                    TN += 1
```

```
total_pairs = TP+TN+FP+FN

expected_index = ((TP+FP)*(TP+FN))/total_pairs
max_index = ((TP+FP)+(TP+FN))/2

# Prevent division by zero

if (max_index-expected_index) == 0:
    return 0.0

ari = (TP-expected_index)/(max_index-expected_index)

return ari

def normalized_mutual_information(labels_true, labels_pred):
    clusters = np.unique(labels_pred)
    classes = np.unique(labels_true)

    matrix = np.zeros((len(classes), len(clusters)))

    for i, c in enumerate(classes):
        for j, k in enumerate(clusters):
            matrix[i, j] = np.sum((labels_true == c) & (labels_pred == k))

    # Entropy calculation
    pi_true = np.sum(matrix, axis=1)/len(labels_true)
    pi_pred = np.sum(matrix, axis=0)/len(labels_pred)

    pi_true = pi_true[pi_true > 0]
    pi_pred = pi_pred[pi_pred > 0]

    h_true = -np.sum(pi_true * np.log2(pi_true))
    h_pred = -np.sum(pi_pred * np.log2(pi_pred))

    mutual_information = 0

    for i in range(len(classes)):
        for j in range(len(clusters)):
            if matrix[i, j] > 0:
                # Probability of point being in class i and cluster j
                p_ij = matrix[i, j] / len(labels_true)
                # Probability of point being in class i
                p_i = np.sum(matrix[i, :]) / len(labels_true)
                # Probability of point being in cluster j
                p_j = np.sum(matrix[:, j]) / len(labels_true)
                mutual_information += p_ij * np.log2(p_ij / (p_i * p_j))
```

```
if (h_true + h_pred) == 0:
    return 0.0

return 2* mutual_information / (h_true + h_pred)

def purity(labels_true, labels_pred):
    clusters = np.unique(labels_pred)
    classes = np.unique(labels_true)

    matrix = np.zeros((len(classes), len(clusters)))

    for i, c in enumerate(classes):
        for j, k in enumerate(clusters):
            matrix[i, j] = np.sum((labels_true == c) & (labels_pred == k))

    dominant_class = np.max(matrix, axis=0)
    return np.sum(dominant_class) / len(labels_true)
```

```
# PCA Class From PCA Code
```

```
class PCA:
    def __init__(self, numberOfPrincipleComponents=2):
        self.numberOfPrincipleComponents = numberOfPrincipleComponents
        self.eigenvectors = None
        self.eigenvalues = None
        self.mean = None
        self.std = None
        self.covarianceMatrix = None
        self.explainedVarianceRatio = None
        self.pca = None
        self.reconstructed = None

    def fit(self, features):
        self.mean = np.mean(features, axis=0)
        self.std = np.std(features, axis=0)

        # Precautionary step to ensure no features having a standard dev
        self.std[self.std == 0] = 1

        standardized = (features - self.mean) / self.std
        self.covarianceMatrix = np.cov(standardized.T)
        eigenvalues, eigenvectors = np.linalg.eig(self.covarianceMatrix)
        # By default, np.argsort sorts in ascending order. For this mod
        sortedIndices = np.argsort(eigenvalues)[::-1]
        self.eigenvalues = eigenvalues[sortedIndices]
```

```
# Sorts columns, not rows.
self.eigenvectors = eigenvectors[:,sortedIndices]

totalVariance = np.sum(eigenvalues)
self.explainedVarianceRatio = self.eigenvalues / totalVariance

return self

def project(self,features):
    standardized = (features - self.mean) / self.std
    # The np.dot function projects the input data onto the provided
    pca = np.dot(standardized, self.eigenvectors[:, :self.numberOfP
    self.pca = pca
    return pca

def reconstruct(self, projectedFeatures):
    # Reconstruct the data after projecting it on the principle axe
    reconstructed = np.dot(projectedFeatures, self.eigenvectors[:, :
    self.reconstructed = reconstructed
    return reconstructed

def reconstructionError(self, features):
    # Calculating the error in the reconstructed data
    projectedData = self.project(features)
    reconstructedData = self.reconstruct(projectedData)

    error = np.mean((features - reconstructedData)**2)

    print(f'Reconstruction Error using MSE:{error:.3f}')
    return error

def plot(self,target):
    # Plotting a scatter plot
    plt.figure()
    color = ['r','g']
    names = ['Benign','Malignant']
    for i in range(len(color)):
        plt.scatter(self.pca[target==i,0], self.pca[target==i,1], c=c
    plt.title("PCA Plot of Two Components")
    plt.xlabel(f"Principal Component 1 With Explained Variance of:{self.explainedVarianceRatio[0]:.3f}")
    plt.ylabel(f"Principal Component 2 With Explained Variance of:{self.explainedVarianceRatio[1]:.3f}")
    plt.legend()

    plt.figure()
    # Plot a bar chart for the explained variance of all features in
    barComponents=np.arange(1, len(self.explainedVarianceRatio)+1)
    plt.bar(barComponents, self.explainedVarianceRatio, align='cente
```

```
plt.ylabel('Explained Variance Ratio')
plt.xlabel('Principle Component Index')
plt.title('Variance Ratio Plot')
```

```
# K-Means Class From K-Means Code

class K_Means:
    def __init__(self, k=2, random_state=None):
        self.k = k
        self.random_state = random_state

    def kmeans_plusplus_init(self, X):
        if self.random_state is not None:
            np.random.seed(self.random_state)

        n_samples, n_features = X.shape
        centers = np.zeros((self.k, n_features))

        # Step 1: Choose first center randomly from data points
        first_center_idx = np.random.randint(0, n_samples)
        centers[0] = X[first_center_idx]
        print(f"Center 1 initialized at index {first_center_idx}")

        # Step 2: Choose remaining k-1 centers
        for i in range(1, self.k):
            # Calculate distance from each point to nearest existing ce
            distances = np.zeros(n_samples)

            for j in range(n_samples):
                # Calculate distances to all existing centers
                dist_to_centers = np.zeros(i)
                for c in range(i):
                    dist_to_centers[c] = np.sum((X[j] - centers[c]) **

            # Store minimum distance (squared distance to nearest c
            distances[j] = np.min(dist_to_centers)

            # Choose next center with probability proportional to dista
            probabilities = distances / np.sum(distances)
            next_center_idx = np.random.choice(n_samples, p=probabiliti
            centers[i] = X[next_center_idx]
            print(f"Center {i+1} initialized at index {next_center_idx}

print(f"\nK-means++ initialization complete!")
print(f"Initial centers shape: {centers.shape}")
```

```
    return centers

def kmeans(self, X, k, centers, epochs, tolerance=1e-6):
    n_samples, n_features = X.shape
    labels = np.zeros(n_samples, dtype=int)
    inertia_history = []

    for epoch in range(epochs):
        # Step 1: Assignment – Assign each point to nearest center
        for i in range(n_samples):
            distances = np.zeros(k)
            for j in range(k):
                distances[j] = np.sum((X[i] - centers[j]) ** 2)
            labels[i] = np.argmin(distances)

        # Step 2: Update – Calculate new centers
        new_centers = np.zeros((k, n_features))
        for j in range(k):
            # Get all points assigned to cluster j
            cluster_points = X[labels == j]
            if len(cluster_points) > 0:
                new_centers[j] = np.mean(cluster_points, axis=0)
            else:
                # If no points assigned, keep old center
                new_centers[j] = centers[j]

        # Calculate inertia (sum of squared distances to nearest ce
        inertia = 0
        for i in range(n_samples):
            inertia += np.sum((X[i] - centers[labels[i]]) ** 2)
        inertia_history.append(inertia)

        # Check for convergence
        center_shift = np.sum(np.abs(new_centers - centers))
        if center_shift < tolerance:
            print(f"\nConverged at epoch {epoch + 1}! (Center shift"
            centers = new_centers
            break

    centers = new_centers
    print(f"Final inertia: {inertia_history[-1]:.4f}")

    # Print cluster sizes
    print(f"\nCluster sizes:")
    for i in range(k):
        cluster_size = np.sum(labels == i)
        print(f"Cluster {i}: {cluster_size} points")
```

```
    return centers, labels, inertia_history

# Starting with experiment

def experiment(features, target, numberOfPrincipleComponents=[2,5,10]):

    results = {
        'dimensions':[], 'mse':[], 'silhouette':[], 'dbi':[], 'chi':[],
        'purity':[], 'wcss':[], 'ari':[], 'explained_variance':[], 'nmi':[]
    }

    classes = len(np.unique(target))

    for n in range(numberOfPrincipleComponents):
        pca = PCA(n)
        pca.fit(features)
        reduced = pca.project(features)

        reconstructed = pca.reconstruct(reduced)
        mse_score = pca.reconstructionError(features)

        kmeans = K_Means(k=classes, random_state=42)
        initialCenters = kmeans.kmeans_plusplus_init(reduced)
        centers, labels, inertia_history = kmeans.kmeans(reduced, classes)

        sil = silhouette_score(reduced, labels)
        dbi = davies_bouldin_index(reduced, labels)
        wcss,chi = calinski_harabasz_index(reduced, labels)
        ari = adjusted_rand_index(target, labels)
        purity_score = purity(target, labels)
        nmi = normalized_mutual_information(target, labels)
        explainedVariance = np.sum(pca.explainedVarianceRatio[:n])

        results['dimensions'].append(n)
        results['mse'].append(mse_score)
        results['silhouette'].append(sil)
        results['dbi'].append(dbi)
        results['chi'].append(chi)
        results['purity'].append(purity_score)
        results['wcss'].append(wcss)
        results['ari'].append(ari)
        results['explained_variance'].append(explainedVariance)
        results['nmi'].append(nmi)

    print(f"\nCurrent Dimensions (Projected Data):{n}\n")
    print(f"MSE:{mse_score}\n")
    print(f"Silhouette Score:{sil}\n")
    print(f"Davies Bouldin Index:{dbi}\n")
```

```
print(f"Calinski Harabasz Index:{chi}\n")
print(f"Purity:{purity_score}\n")
print(f"WCSS:{wcss}\n")
print(f"ARI:{ari}\n")
print(f"Explained Variance:{explainedVariance}\n")
print(f"NMI:{nmi}\n")

if n==2:
    print("\nConfusion Matrix\n")

    cm = confusion_matrix(target, labels)
    plt.figure()
    sb.heatmap(cm, annot=True, fmt='g', cbar=False, xticklabels=['C'
    plt.xlabel('Predicted Cluster ID')
    plt.ylabel('True Label')
    plt.title('Confusion Matrix')
    plt.show()

    print("\n2D Projection and Variance Plot")
    pca.plot(target)
    plt.show()

return results

def plotWCSS(results):
    plt.figure()
    plt.title("Elbow Curve")
    plt.xlabel("Number of dimensions")
    plt.ylabel("WCSS")
    plt.plot(results['dimensions'], results['wcss'], label='WCSS')
    plt.legend()
    plt.grid(True)
    plt.show()

def plotInternal(results):
    plt.figure()
    plt.title("Internal Validation Metrics")
    plt.xlabel("Number of dimensions")
    plt.ylabel("Score")
    plt.plot(results['dimensions'], results['silhouette'], label='Sil
    plt.plot(results['dimensions'], results['dbi'], label='Davies Bou
    plt.legend()
    plt.grid(True)
    plt.show()

def plotCHI(results):
    plt.figure()
```

```
plt.title("Calinski Harabasz Index")
plt.xlabel("Number of dimensions")
plt.ylabel("Score")
plt.plot(results['dimensions'], results['chi'], label='Calinski H
plt.legend()
plt.grid(True)
plt.show()

def plotExternal(results):
    plt.figure()
    plt.title("External Validation Metrics")
    plt.xlabel("Number of dimensions")
    plt.ylabel("Score")
    plt.plot(results['dimensions'], results['ari'], label='ARI')
    plt.plot(results['dimensions'], results['purity'], label='Purity')
    plt.plot(results['dimensions'], results['nmi'], label='NMI')
    plt.legend()
    plt.grid(True)
    plt.show()

def plotReconstructionError(results):
    plt.figure()
    plt.title("Reconstruction Error")
    plt.xlabel("Number of dimensions")
    plt.ylabel("Score")
    plt.plot(results['dimensions'], results['mse'], label='MSE')
    plt.legend()
    plt.grid(True)
    plt.show()
```

```
print("Experiment 3: K-Means After PCA\n")
results = experiment(features, target, number_of_principle_components=2)

print("\nVisualization\n")
plot_WCSS(results)
plot_Reconstruction_Error(results)

plot_Internal(results)
plot_CHI(results)
plot_External(results)
```

Experiment 3: K-Means After PCA

Reconstruction Error using MSE: 809.793

Center 1 initialized at index 102

Center 2 initialized at index 430

```
K-means++ initialization complete!
Initial centers shape: (2, 2)

Converged at epoch 5! (Center shift 0.000000 < tolerance 1e-06)
Final inertia: 5332.5556

Cluster sizes:
Cluster 0: 378 points
Cluster 1: 191 points

Current Dimensions (Projected Data):2

MSE:809.7927504717508

Silhouette Score:0.5084690190672025

Davies Bouldin Index:0.8467403809951899

Calinski Harabasz Index:293705.1411818269

Purity:0.9068541300527241


WCSS:5332.555574203135

ARI:0.6592305570949524

Explained Variance:0.6324320765155946


NMI:0.5403814627831337
```

Confusion Matrix







2D Projection and Variance Plot



Variance Ratio Plot





Reconstruction Error using MSE:729.817

Center 1 initialized at index 102

Center 2 initialized at index 430

K-means++ initialization complete!

Initial centers shape: (2, 5)

Converged at epoch 7! (Center shift 0.000000 < tolerance 1e-06)

Final inertia: 8992.1305

Cluster sizes:

Cluster 0: 375 points

Cluster 1: 194 points

Current Dimensions (Projected Data):5

MSE:729.8173415293901

Silhouette Score:0.3918571464102746

Davies Bouldin Index:1.1489013654222102

Calinski Harabasz Index:241018.71792267804

Purity:0.9050966608084359

WCSS:8992.130462630234

ARI:0.6536246043910179

Explained Variance:0.8473427431680726

NMI:0.5324078598532429

Reconstruction Error using MSE:515.558

Center 1 initialized at index 102

Center 2 initialized at index 432

K-means++ initialization complete!

Initial centers shape: (2, 10)

Converged at epoch 4! (Center shift 0.000000 < tolerance 1e-06)

Final inertia: 10770.5470

Cluster sizes:

Cluster 0: 375 points

Cluster 1: 194 points

```
Current Dimensions (Projected Data):10
MSE:515.5581517569499
Silhouette Score:0.3564607021923901
Davies Bouldin Index:1.2679003700826144
Calinski Harabasz Index:227728.91900205094
Purity:0.9050966608084359
WCSS:10770.547027720138
ARI:0.6536246043910179
Explained Variance:0.9515688143366668
NMI:0.5324078598532429
Reconstruction Error using MSE:249.052
Center 1 initialized at index 102
Center 2 initialized at index 432
K-means++ initialization complete!
Initial centers shape: (2, 15)
Converged at epoch 3! (Center shift 0.000000 < tolerance 1e-06)
Final inertia: 11365.5705
Cluster sizes:
Cluster 0: 374 points
Cluster 1: 195 points
Current Dimensions (Projected Data):15
MSE:249.05152377929303
Silhouette Score:0.3460830257023994
Davies Bouldin Index:1.3064723052258922
Calinski Harabasz Index:223994.9106473569
Purity:0.9068541300527241
WCSS:11365.570468457017
ARI:0.6594271719206948
Explained Variance:0.9864881227145469
NMI:0.5380305878158959
```

```
Reconstruction Error using MSE:51.404
Center 1 initialized at index 102
Center 2 initialized at index 432
```

```
K-means++ initialization complete!
Initial centers shape: (2, 20)
```

```
Converged at epoch 3! (Center shift 0.000000 < tolerance 1e-06)
Final inertia: 11520.1659
```

```
Cluster sizes:
Cluster 0: 374 points
Cluster 1: 195 points
```

```
Current Dimensions (Projected Data):20
```

```
MSE:51.40359070901072
```

```
Silhouette Score:0.3438371031140346
```

```
Davies Bouldin Index:1.3159467193532868
```

```
Calinski Harabasz Index:223037.94352372066
```

```
Purity:0.9068541300527241
```


```
WCSS:11520.165852974762
```




```
ARI:0.6594271719206948
```


```
Explained Variance:0.9955720432628138
```

```
NMI:0.5380305878158959
```


Visulization









Calinski Harabasz Index



External Validation Metrics





## The Curse of Dimensionality

The data was seen to be highly dependent on a number of features. This lead the reduction in dimensions to lead to fascinating results when combined with the K-Means Algorithm.

The results prove that using two dimensions is the most optimal number of dimensions to be chosen. The silhouette score was 0.508 when the number of dimensions chosen was 2, and eventually dropped down to 0.344 when they increased to 20. Adding more dimensions introduced noise to the features, increasing the distances between points of the same cluster and degrading the density of each cluster.

As would be expected, the reconstruction error significantly decreases as the number of dimensions increase (809.8 to 51.4) yet the adjusted rand index did not change and was almost 0.6. This goes on to show that the first two components were helpful in identifying almost 60% of the variance in the data and was a 100% useful for clustering. Any increase in dimensions did not yield significant results for the separation between the two clusters.

